

A Faithful Lean Account of the PZG–BEDC Kernel

Unified Theory Working Line

Abstract

We report the portion of the PZG–BEDC dynamic ontology kernel formalized in Lean 4 with `mathlib`. PZG denotes prime-axis Zeckendorf coding: each prime axis carries a finite Zeckendorf word, and decoding first turns those words into prime exponents and then reconstructs a positive natural number. The checked development covers 86 content modules plus the aggregation module: the finite generation layer, Newman normal forms, prime-axis factorization and Euclid escape, Zeckendorf and PZG bijections, finite golden-ratio algebra, deficit and golden-jump ledgers, golden double-sided and fiber-coordinate identities, Cassini–Fricke trace conservation, finite readings and two-square classifiers, Kleene/Cantor/Gödel barriers for self-coding, cost and hidden-fiber rigidity, critical-line reflection algebra, Euler products, golden-weight finite Euler factorization, Weil explicit-formula small-support prime-side reduction, concrete archimedean well-definedness, the heart-electrode residual-identity foundation, the Laplace-transform convergence abscissa and half-plane holomorphy that form the analytic base of the Landau oscillation engine, the golden coordinate ring $\mathbb{Z}[\varphi]$ with its norm, conjugation automorphism and units (the φ -coordinate system L0), no-zero-line zeta facts, completed zeta functional equations, and the RH ledger equivalence. The generated-results section records twenty-nine machine-checked additions; together with the eight checked groups folded into the formalized-results section, the manuscript records thirty-seven such additions, including the unconditional shifted Zeckendorf–Beatty bridge, the golden double-sided/fiber identities, the heart-electrode foundation, the Laplace-transform abscissa and half-plane holomorphy, and fourteen oracle-suggested conjectures certified locally after proposal. Every mathematical assertion labeled as machine checked below names an actual Lean declaration under `papers/unified_theory/lean4/UnifiedTheory`. The RH result is an equivalence with `mathlib`’s `RiemannHypothesis`, not a proof of RH; quasicrystal zeta, full explicit-formula zero-side analysis, Weil archimedean positivity, unconditional small-support positivity, the analytic golden-weight Dirichlet L-series, positivity bridges, the empirical heart Sturmian shell, and open leaves such as O-5 and O-6 remain outside the checked-result column. The Lean project is `mathlib`-based and isolated from the `mathlib`-free BEDC core; the stated reproducibility surface is the branch `feat/unified-theory-mathlib`, with the full project check given by `lake build`.

1 Introduction

The source theory proposes a dynamic mathematical ontology kernel

$$\mathcal{K}_{\text{PZG}}$$

whose objects are not only static data, but generated histories, standard codes, decoders, normalizers, finite readings, and ledger states. The PZG part of the name is literal: *prime axes* provide the coordinate directions, while *Zeckendorf words* provide the per-axis digit system. In the formalized carrier, a PZG table is a finite-support map

$$p \longmapsto w_p$$

where the support is restricted to primes and each w_p is a legal Zeckendorf representation. Decoding maps each w_p to a natural exponent and then uses unique factorization to recover an element of $\mathbb{N}_{>0}$.

The goal of this Lean line is not to certify the full speculative reach of the source document. It is narrower and more valuable as a logical audit: identify the finite and algebraic claims that can be stated cleanly, prove them in Lean, and keep analytic or self-referential frontier claims outside the checked-result column until they have genuine declarations and proofs. This paper therefore distinguishes three classes.

- (i) **Checked theorem or definition:** a declaration in `UnifiedTheory/**/*.lean` that can be referenced by name.
- (ii) **Encoded statement:** a Lean `Prop` used to record a target without pretending that it is proved.
- (iii) **Outlook:** mathematical plans or analytic programs that are not yet formalized.

The development is deliberately `mathlib`-based. It uses `mathlib`'s APIs for natural-number factorization, Zeckendorf representations, golden-ratio identities, pigeonhole arguments, and Fermat's theorem on sums of two squares. This is separate from the `mathlib`-free BEDC core in the main repository: no file under the BEDC core imports this Lean project, and this project does not alter the BEDC no-`mathlib` invariant.

2 Formalized Results

This section lists the checked content by layer. Each item gives an informal statement, the precise Lean declaration name, and the reason it captures the corresponding part of the source theory.

2.1 Generation, Mark Exchange, and Histories

Two marks and exchange. The finite basis of the generation layer is the inductive type

`UnifiedTheory.Mark`.

It has exactly two constructors, `zero` and `one`. The exchange operation is

`UnifiedTheory.Mark.sigma : Mark -> Mark`,

with checked facts

`UnifiedTheory.Mark.sigma_involutive`

and

`UnifiedTheory.Mark.sigma_nontrivial`.

Informally, the first theorem says $\sigma(\sigma(x)) = x$ for both marks, and the second records that σ really swaps the two constructors. This is the formal core of the source definitions 1.1–1.2: the minimal binary marker and its nontrivial exchange are not axioms, but a two-constructor inductive type and a case split.

Finite histories. The history carrier is

`UnifiedTheory.MarkHist`,

with constructors `empty` and `ext`. The formal append operation satisfies

`UnifiedTheory.MarkHist.empty_append`,

`UnifiedTheory.MarkHist.append_empty`,

`UnifiedTheory.MarkHist.append_assoc`,

and its length satisfies

`UnifiedTheory.MarkHist.length_append.`

These are the checked version of the finite-history monoid laws in proposition 2.4. The source document describes generation induction and event finiteness as axiomatic principles; the Lean line does not add such axioms. It realizes the finite layer by ordinary inductive recursion and lists.

Events and generation. Events are structures

`UnifiedTheory.Event`

with fields for source history, opcode, argument, and output tag. Event histories are lists, and generation is list extension:

`UnifiedTheory.EventHist.generate.`

The checked statement

`UnifiedTheory.EventHist.generate_length`

says that appending one event increases the list length by one. This is the finite operational shadow of the source idea that a generated event extends a history.

2.2 Unicity by Abstract Rewriting

The formal rewriting module is generic in a type α and a relation $r : \alpha \rightarrow \alpha \rightarrow \text{Prop}$. It defines the reflexive-transitive closure

`UnifiedTheory.Rewriting.Star,`

normal forms

`UnifiedTheory.Rewriting.Normal,`

joinability, local confluence, confluence, and termination:

`UnifiedTheory.Rewriting.Joinable,`

`UnifiedTheory.Rewriting.LocallyConfluent,`

`UnifiedTheory.Rewriting.Confluent,`

`UnifiedTheory.Rewriting.Terminates.`

The checked Newman theorem is

`UnifiedTheory.Rewriting.newman_confluent.`

Its type is:

$\text{Terminates}(r) \rightarrow \text{LocallyConfluent}(r) \rightarrow \text{Confluent}(r).$

The normal-form theorem is

`UnifiedTheory.Rewriting.exists_unique_normal_form,`

which states that under the same hypotheses every element has a unique normal form. This exactly formalizes the source theorem 3.2 at the abstract rewriting level. It does not assert that every future normalizer in the theory has already been given a terminating rewrite system; it supplies the reusable criterion.

2.3 Prime Axes and Euclid Escape

The prime-axis state space is

`UnifiedTheory.PrimeExp,`

a finite-support natural exponent vector whose support consists of primes. The declaration

`UnifiedTheory.PrimeExp.equivPNat`

has type

$\mathbb{N}_{>0} \simeq \text{UnifiedTheory.PrimeExp}.$

This is the formalized prime-axis statement: positive natural numbers are equivalent to finite prime exponent vectors. The proof is a mathlib wrapper around `Nat.factorizationEquiv`, so the paper should not claim that this Lean file rederives the full elementary proof of unique factorization from scratch.

The exponent reading is

`UnifiedTheory.vp,`

and the multiplicative additivity theorem is

`UnifiedTheory.vp_mul.`

It says

$$v_p(mn) = v_p(m) + v_p(n)$$

for $m, n \in \mathbb{N}_{>0}$. This is the checked content behind theorem 4.4’s “multiplication is coordinate-wise addition on prime axes” slogan.

The group-ification of this free structure to the positive rationals (source theorem 18.10) is checked on its free half: `UnifiedTheory.primes_mul_indep` proves that for a formal integer exponent vector $a : \mathbb{N} \rightarrow_0 \mathbb{Z}$ with prime support, $\prod_p p^{a_p} = 1$ in $\mathbb{Q}$ forces $a = 0$. Equivalently the map $a \mapsto \prod_p p^{a_p}$ is injective, so the primes are multiplicatively independent and $(\mathbb{Q}_{>0}, \times)$ is free (no relations) on the prime axes. The proof applies the p -adic valuation v_q to the product, which splits as a sum of valuations (`UnifiedTheory.padicValRat_prod`); distinct-prime terms vanish (`UnifiedTheory.padicValRat_distinct_prime`), the $p = q$ term returns a_q via the integer-power valuation (`UnifiedTheory.padicValRat_zpow`), and $v_q(1) = 0$ forces $a_q = 0$. The surjective half is also checked: `UnifiedTheory.primes_generate` proves that every positive rational is such a product, taking the exponent vector $a = v(\text{num}) - v(\text{den})$ from the factorizations of the coprime numerator and denominator and splitting the product by `Finsupp.prod_add_index` into $\text{num}/\text{den} = q$. Together these two halves are source theorem 18.10: $(\mathbb{Q}_{>0}, \times)$ is the free abelian group on the prime axes, the group-ification of theorem 4.4.

Euclid escape is represented by

`UnifiedTheory.euclidEscape`

and proved by

`UnifiedTheory.exists_prime_not_mem.`

For every finite set S of primes, the theorem produces a prime divisor of

$$\prod_{p \in S} p + 1$$

that is not in S . This is the formalized version of theorem 4.5: no finite window of prime axes exhausts the prime directions.

2.4 Zeckendorf Words and the PZG Bijection

On a single axis, the Zeckendorf carrier is

`UnifiedTheory.AxisWord,`

a subtype of lists satisfying mathlib's `List.IsZeckendorfRep`. The equivalence

`UnifiedTheory.AxisWord.equivNat`

has type

$\mathbb{N} \simeq \text{UnifiedTheory.AxisWord}.$

The encode-decode inverse theorems

`UnifiedTheory.AxisWord.decode_encode, UnifiedTheory.AxisWord.encode_decode`

state the existence and uniqueness directions of Zeckendorf representation. This is the checked theorem 5.3, again using mathlib's Zeckendorf API rather than a handwritten proof.

The PZG table carrier is

`UnifiedTheory.PZGTable.`

It is not merely an exponent vector: it keeps the per-prime Zeckendorf word. The axis-level equivalence

`UnifiedTheory.PZGTable.equivPrimeExp`

decodes each axis word to an exponent and preserves prime support. Composing this with prime factorization gives

`UnifiedTheory.PZGTable.decode : PZGTable \simeq \Npos .`

The checked inverse laws

`UnifiedTheory.PZGTable.decode_encode, UnifiedTheory.PZGTable.encode_decode`

are the formal statement of theorem 5.5: PZG tables and positive natural numbers are in bijection.

Normalization is defined canonically:

`UnifiedTheory.PZGTable.normAdd z w = encode (decode z * decode w).`

The theorem

`UnifiedTheory.PZGTable.decode_normAdd`

says

$\text{decode}(\text{normAdd}(z, w)) = \text{decode}(z) \text{ decode}(w),$

and

`UnifiedTheory.PZGTable.vp_decode_normAdd`

says the same thing coordinatewise as prime-exponent addition. This captures theorem 5.6 as a canonical normalization theorem. It is not a proof of a step-by-step carry algorithm or carry termination; it defines the normalized result by the uniqueness of the PZG bijection.

2.5 Golden Double-Sided Algebra

The golden algebra carrier is

`UnifiedTheory.PhiInt.`

An element is an integer pair representing $a + b\varphi$ with $\varphi^2 = \varphi + 1$. Multiplication, conjugation, and real embeddings are defined by

`UnifiedTheory.PhiInt.conj,`
`UnifiedTheory.PhiInt.toRealPlus,`
`UnifiedTheory.PhiInt.toRealMinus.`

The theorem

`UnifiedTheory.PhiInt.fixed_conj_iff`

says that a `PhiInt` is fixed by conjugation exactly when its φ -coefficient is zero. This is the algebraic reason the deficit integrality proof works.

The embedding compatibility theorem

`UnifiedTheory.PhiInt.toRealMinus_eq_toRealPlus_conj`

identifies the shrinking real face with the expanding real face after conjugation, while

`UnifiedTheory.PhiInt.toRealPlus_mul`

states that the expanding embedding is multiplicative. The Fibonacci decomposition of powers is

`UnifiedTheory.PhiInt.phiPow_ab,`

which says that the b -coordinate of φ^n is the Fibonacci number F_n .

The kernel's cascade cancellation rests on a one-page golden algebra, checked in Lean. From $\varphi^2 = \varphi + 1$ one gets $\varphi^3 = 2\varphi + 1$, $\varphi^4 = 3\varphi + 2$, the identity $\varphi^4 + \varphi^2 - 1 = 2\varphi^3$, `UnifiedTheory.goldenRatio_pow4_add_sq_sub_one`, the second-order golden cancellation $\varphi^2 + \varphi^3 = \varphi^4$, `UnifiedTheory.goldenRatio_sq_add_cube`, and the Zeckendorf $11 \rightarrow 100$ analytic shadow $\varphi^4 - \varphi^3 - \varphi^2 = 0$, `UnifiedTheory.goldenRatio_pow4_sub_cube_sub_sq`. The conjugate satisfies $\psi = -1/\varphi$, `UnifiedTheory.goldenConj_eq_neg_inv`, and the constant C_0 splits into two half-goldens $\varphi/2 = \sqrt{5}/2 - 1/(2\varphi)$, `UnifiedTheory.goldenRatio_half_eq_sqrt5_half_sub`. The one-line Binet identity seeding the $\bar{r}$ -chain, $F_{k+1}/\varphi = F_k - \psi^{k+1}$, is `UnifiedTheory.fib_succ_div_goldenRatio`.

The two-dimensional lift is forced at the outset: `UnifiedTheory.no_geometric_fib_ratio` proves there is no constant ratio Λ with $F_{k+2} = \Lambda F_{k+1}$ for all k (the consecutive-Fibonacci ratios $1, 2, 3/2, \dots$ are non-constant), so no single geometric weight reproduces the Zeckendorf decode — this is source theorem 6.2, the structural no-go that motivates the double-sided eigenline lift. The double-sided scalar identities are then explicit. The constant c^* from source 6.145 is `UnifiedTheory.cstar_eq_sqrt5_phi`, namely $\varphi^2 + 1 = \sqrt{5}\varphi$; Binet's difference formula `UnifiedTheory.goldenPow_sub_goldenConjPow` proves $\varphi^n - \psi^n = \sqrt{5}F_n$, so the source definition 6.4 length $\ell(z) = (\lambda_+ - \lambda_-)/\sqrt{5}$ is an integer Fibonacci-weighted quantity, not an irrational golden-weighted one. The product formula `UnifiedTheory.goldenPow_mul_goldenConjPow`, $\varphi^n \psi^n = (-1)^n$, is the double-sided orientation sign of the alternating $\bar{r}$ -chain. The sum completes the triple of symmetric functions of the golden-rotation eigenvalues: `UnifiedTheory.goldenPow_add_golden` proves the integer Lucas trace $\varphi^n + \psi^n = 2F_{n+1} - F_n$. Finally, `UnifiedTheory.floorPhiMul_eq` proves $\lfloor (v+1)\varphi \rfloor = (v+1) + \lfloor (v+1)/\varphi \rfloor$, and `UnifiedTheory.beattyShiftRead_eq_add_floor` proves $S(v) = v + \lfloor (v+1)/\varphi \rfloor$; hence shift decode and the fiber coordinate $a + b = \lfloor (v+1)/\varphi \rfloor$ are floors of one golden rotation, as in source Corollary 6.48'.

For axis words, the double-sided reading is

`UnifiedTheory.AxisWord.betaPhi.`

The checked gap identity

`UnifiedTheory.AxisWord.betaPhi_gap`

says that the difference between the expanding and shrinking real readings is $\sqrt{5}$ times the ordinary decoded value. At the PZG-table level, the two real readings are

`UnifiedTheory.PZGTable.lambdaPlus,` `UnifiedTheory.PZGTable.lambdaMinus.`

These declarations formalize definition 6.4 as a finite sum over the finite support of a PZG table.

Carry ledger identities. The finite golden carry identities are checked as

`UnifiedTheory.Golden.gold_carry_plus,` `UnifiedTheory.Golden.gold_carry_minus,`

and the bottom-rule identities as

`UnifiedTheory.Golden.gold_bottom_plus,` `UnifiedTheory.Golden.gold_bottom_minus.`

The exact `PhiInt` version is

`UnifiedTheory.Golden.phiPow_carry.`

These are the formal finite algebra behind proposition 6.21: adjacent Fibonacci carries preserve the double-sided account, while the bottom identity records the unit deficit equation $2\varphi^2 - \varphi^3 = 1$.

Deficit integrality. The normalization deficit is defined by

`UnifiedTheory.deficitPhi.`

The theorem

`UnifiedTheory.deficitPhi_integer`

states that for all $v, w : \mathbb{N}$ there exists $k : \mathbb{Z}$ such that

$$\text{deficitPhi}(v, w) = \text{ofInt}(k).$$

The equivalent conjugation-fixed statement is

`UnifiedTheory.deficitPhi_conj_fixed.`

This is the checked theorem 6.22: the deficit is an integer because its φ -coordinate cancels.

Deficit trichotomy and the shrinking window. The stronger three-value statement is checked unconditionally. The theorem

`UnifiedTheory.abs_betaMinus_le`

proves the shrinking-face window bound

$$|\beta_-(\text{encode}(n))| \leq 1/\varphi$$

for every natural number n . The proof uses the gap condition in Zeckendorf words and a finite geometric estimate

`UnifiedTheory.geomBound.`

Combined with deficit integrality, this gives

`UnifiedTheory.deficitTrichotomy_holds,`

the checked theorem 6.25: the integer deficit lies in

$$\{-1, 0, 1\}.$$

This is the first window-length bearing theorem in the Lean development and proves the target

`UnifiedTheory.DeficitTrichotomy`

into a proved theorem.

Beatty-floor deficit and carry conservation. The Beatty deficit module gives an independent floor-arithmetic route to the same three-value phenomenon. It defines the shifted Beatty reading

`UnifiedTheory.S`

and the binary carry deficit

`UnifiedTheory.cDef.`

The theorem

`UnifiedTheory.cDef_mem`

proves directly that `cDef a b` is always -1 , 0 , or 1 . The checked cocycle identity

`UnifiedTheory.carry_conservation`

states

$$c(a, b) + c(a + b, c) = c(a, b + c) + c(b, c),$$

which is the formal carry-conservation law 6.62, or $\delta^2 = 0$, for the Beatty deficit ledger. The bridge from this Beatty reading to shifted Zeckendorf sums is now closed:

`UnifiedTheory.S_eq_shifted_zeck`

is now an unconditional theorem. The asymmetric shrinking-window bound

`UnifiedTheory.minusWindow_holds`

is proved without hypotheses by splitting the signed Zeckendorf chain into its parity components and applying the strict geometric estimate

`UnifiedTheory.geomBoundStrict`

together with the signed window theorem

`UnifiedTheory.signedGeom_window.`

The algebraic bridge

`UnifiedTheory.betaMinus_encode_eq`

identifies the Beatty deficit with the shifted Zeckendorf sum, and

`UnifiedTheory.shiftedZeckendorf_of_minusWindow`

turns the window bound into the bridge. Thus theorem 6.44 is checked as the unconditional identity

$$S(v) = \sum_{i \in \text{encode } v} \text{fib}(i + 1).$$

n -ary deficit. The list-valued Beatty deficit is

`UnifiedTheory.Cn.`

The decomposition theorem

`UnifiedTheory.Cn_cons`

says that adjoining an entry decomposes the list deficit into one binary deficit plus the tail deficit. The checked trapezoid estimate

`UnifiedTheory.Cn_abs_le`

states

$$|C_n(v_1, \dots, v_m)| \leq m - 1.$$

This is the symmetric part of theorem 6.60. The sharper asymmetric endpoint count is not claimed here.

Finite Fibonacci core. The finite recurrence

`UnifiedTheory.Golden.finiteW_recurrence`

and Cassini identity

`UnifiedTheory.Golden.cassini`

formalize the finite algebraic core behind the trace-map and generating-function discussion. The traced logarithmic coordinate

`UnifiedTheory.uCoord`

and quadratic form

`UnifiedTheory.qForm`

support the Cassini–Fricke conservation law

`UnifiedTheory.cassini_fricke.`

It states that

$$Q(u_{K+1}, u_K) = 5xy(-1)^{K+1}$$

for

$$u_K = -x\varphi^{K+1} + y\psi^{K+1}.$$

This is the checked theorem 6.36 at the finite trace-map level. These theorems do not formalize the infinite function $W(x, y)$, the quasicrystal zeta object, its Euler-product or continuation theory, or zero transport.

2.6 Golden-Weight Generating Function and Finite Euler Product

The golden weight Ω_φ is additive over coprime factors, so its generating function $n \mapsto x^{\Omega_\varphi(n)}$ is a multiplicative arithmetic function `UnifiedTheory.goldGenAF`, checked in `UnifiedTheory.goldGenAF_isMultiplicative`. Its local factor at a prime power is $x^{S(v)}$, because $\Omega_\varphi(p^v) = S(v)$:

`UnifiedTheory.goldWeight_prime_pow`

`UnifiedTheory.goldGenAF_prime_pow.`

Consequently, for every $n \neq 0$ the generating function factors as the finite Euler product

$$x^{\Omega_\varphi(n)} = \prod_{p^k \parallel n} x^{S(k)},$$

checked as `UnifiedTheory.goldGenAF_euler`. This is the algebraic Euler factorization of the golden weight over prime axes. The golden weight is additive over coprime factors but not completely additive ($S(2) = 3 \neq 2S(1) = 4$), so the generating function is multiplicative yet *not* completely multiplicative: at $x = 2$, `goldGenAF 2 (2 · 2) = 2S(2) = 8` while `(goldGenAF 2 2)2 = (2S(1))2 = 16`, checked as `UnifiedTheory.goldGenAF_not_completelyMultiplicative`. Hence the local Euler factor is genuinely non-geometric and the golden L-series does not reduce to a shift or power of ζ . The analytic Dirichlet series $\sum_n x^{\Omega_\varphi(n)} n^{-s}$ —convergence on a half-plane, meromorphic continuation, and any relation to ζ —is not formalized and remains an open leaf.

2.7 Golden Weight, Squarefreeness, and Primality

Because $S(v) \geq 2$ for $v \geq 1$ (with $S(1) = 2$ and S monotone), the golden weight is bounded below by twice the number of distinct prime factors, $2\omega(n) \leq \Omega_\varphi(n)$, checked as `UnifiedTheory.two_mul_primeFactors`. Equality holds exactly for squarefree n , since $\Omega_\varphi(n) = 2\omega(n)$ forces every prime exponent to equal 1:

$$\Omega_\varphi(n) = 2\omega(n) \iff n \text{ squarefree},$$

checked as `UnifiedTheory.goldWeight_eq_two_mul_primeFactors_card_iff_squarefree`. On the prime-power axis the golden weight reads primality: for n a prime power, $\Omega_\varphi(n) = 2$ iff n is prime, checked as `UnifiedTheory.isPrimePow_goldWeight_eq_two_iff_prime` (as $S(v) = 2 \iff v = 1$). Thus weight 2 is the first-order prime channel of the golden weight.

The next weight value detects prime squares: for nonzero n , $\Omega_\varphi(n) = 3$ iff $n = p^2$ for some prime p , checked as `UnifiedTheory.goldWeight_eq_three_iff_prime_sq`. The bound $2\omega(n) \leq \Omega_\varphi(n)$ forces $\omega(n) = 1$ (a prime power), and then $S(v) = 3 \iff v = 2$ pins the exponent to 2; this rules out semiprimes ($\omega = 2$, so $\Omega_\varphi \geq 4$), cubes ($S(3) = 5$), and non-prime squares. The golden weight also evaluates cleanly on factorials: by Legendre's factorial factorization,

$$\Omega_\varphi(N!) = \sum_{p \mid N!} S\left(\sum_{i \geq 1} \lfloor N/p^i \rfloor\right),$$

checked as `UnifiedTheory.goldWeight_factorial`, the inner sum being the p -adic valuation of $N!$.

2.8 Finite Readings and Two-Square Classification

The finite-reading pigeonhole theorem is

`UnifiedTheory.Reading.reading_has_fiber.`

Given a finite carrier, a finite class set, and a map into that class set, if the class set has smaller cardinality than the carrier, two distinct elements share the same reading. This proves theorem 8.5 at the level of finite sets: a finite reading with too few values has a nontrivial fiber.

The two-square classifier includes the prime-level theorem

`UnifiedTheory.Reading.prime_sq_add_sq_iff.`

For a prime p , it states

$$(\exists a b : \mathbb{N}, a^2 + b^2 = p) \iff p \bmod 4 \neq 3.$$

The easy obstruction is proved in the file from the fact that squares modulo 4 are 0 or 1, and the existence direction uses `mathlib`'s theorem for primes that are sums of two squares. This captures the prime-classifier part of chapter 9.

The general classifier is checked:

`UnifiedTheory.Reading.sq_add_sq_iff_factorization.`

It states that a natural number n is a sum of two squares exactly when every prime factor $p \equiv 3 \pmod{4}$ appears in the factorization of n with even exponent. This is theorem 9.10 in its ordinary arithmetic form. The proof is a `mathlib`-based wrapper around the standard sum-of-two-squares factorization theorem; the Lean file records the kernel-facing reading theorem, not an elementary reproof of Fermat's theorem.

The constructive input behind the $p \equiv 1 \pmod{4}$ branch is source theorem 9.5, checked here as an *explicit* certificate rather than a mere existence claim. The factorial pairing `UnifiedTheory.factorial_pair`

proves $(p-1)! \equiv (-1)^m(m!)^2$ in $\mathbb{Z}/p$ for odd prime p and $m = (p-1)/2$, by reflecting the upper half $\{m+1, \dots, p-1\}$ onto the lower half via $k \mapsto p-k$ and using $\overline{p-k} = -\bar{k}$ to extract $(-1)^m$. Combined with Wilson’s lemma and m even for $p \equiv 1 \pmod{4}$, this yields `UnifiedTheory.neg_one_sq_witness`: the element $m!$ is an *explicit* square root of -1 modulo p . This strengthens `mathlib`’s existence-only `ZMod.exists_sq_eq_neg_one_iff` to a named witness, in keeping with the kernel’s certificate discipline.

2.9 Self-Coding and Diagonal Barriers

The self-code layer contains Cantor diagonal theorems for arbitrary types. The declaration

`UnifiedTheory.SelfCode.no_classifier_surjective`

says that no classifier

$$X \rightarrow \text{Set}(X)$$

is surjective. The declaration

`UnifiedTheory.SelfCode.diagonal_not_classified`

exhibits the diagonal property $\{y \mid y \notin \text{classify}(y)\}$ and proves it is not hit by any index. The dual declaration

`UnifiedTheory.SelfCode.no_decoder_injective`

says that no decoder $\text{Set}(X) \rightarrow X$ is injective.

These are theorem 16.2–16.3 in their pure diagonal form.

The same layer contains the Kleene fixed-point theorem in `mathlib`’s partial-recursive-code semantics:

`UnifiedTheory.SelfCode.kleene_fixed_point.`

For every computable transformation on `Nat.Partrec.Code`, it produces a code whose partial-function behavior agrees with the transformed code. The bundled chapter statement

`UnifiedTheory.SelfCode.sc1_chapter16_unconditional`

combines Kleene fixed points with the Cantor classifier and decoder barriers. This discharges the machine axiom U at the self-code level by moving to partial-function semantics: no total-termination assumption is needed for theorem 16.1’. It still does not formalize the later internalization tower, same-layer closure predicate, verified substitution evaluator, or Gödel-style arithmeticized residual.

The arithmetic sequence encoder

`UnifiedTheory.SelfCode.godelEncode`

maps a finite list $(l_0, \dots, l_{m-1})$ to the prime product $\prod_i p_i^{l_i+1}$. Its injectivity theorem

`UnifiedTheory.SelfCode.godelEncode_injective`

is theorem 14.1: two encoded lists with the same Gödel product are equal. This supplies the checked arithmetic infrastructure for self-code separation; the full syntax tower and substitution evaluator remain outside the checked column.

For the fixed-length Gödel-vector code $\text{Gvec}(x) = \prod_i p_i^{x_i+1}$, perfect-power collapse is coordinatewise. The checked criterion

`UnifiedTheory.SelfCode.gvec_isPow_iff_coord_dvd`

says that $\text{Gvec}(x)$ is a perfect e -th power exactly when every shifted coordinate exponent $x_i + 1$ is divisible by e ; in particular it is a square exactly when every $x_i + 1$ is even. The golden weight of the same code is read by

$$\Omega_\varphi(\text{Gvec}(x)) = \sum_i S(x_i + 1),$$

and

`UnifiedTheory.SelfCode.gvec_Omega_first_shell`

identifies the first non-base golden shell: $\Omega_\varphi(\text{Gvec}(x)) = 2k + 1$ exactly when some $x_i = 1$ and every other $x_j = 0$. These are elementary finite-code facts: they rule out hidden perfect-power collapse of the code and pin the first golden shell to single unit perturbations.

2.10 Dynamics and the Cost Arrow

For event histories, the checked cost theorem is

`UnifiedTheory.EventHist.cost_lt_generate.`

It states that generating one event strictly increases the length of the event history. The theorem

`UnifiedTheory.EventHist.generate_ne`

then says that the generated history is not equal to the original history.

This is the formal content of theorem 18.11–18.12 represented in Lean: irreversibility is expressed as strict growth of the event ledger. It is not the full prime-address dynamics with logarithmic cost over signed exponent states.

The prime-axis logarithmic length is formalized as

`UnifiedTheory.L.`

The checked positivity theorem

`UnifiedTheory.L_pos`

says that every nonzero prime-exponent state has strictly positive logarithmic length. This is theorem 18.7 for the prime-axis state carrier. It supplies the length input for the phase and critical-line ledger, but it does not by itself define a completed heat trace or analytic continuation.

The Fibonacci–Fricke trace map $T(a, b, c) = (b, c, bc - a)$ also has a finite quotient-dynamics constraint under the three even-sign flips $(a, b, c) \mapsto (-a, -b, c)$, $(-a, b, -c)$, and $(a, -b, -c)$. It does not commute with the individual flips; instead it cyclically permutes them:

$$T \circ \text{flipAB} = \text{flipAC} \circ T$$

and cyclically, checked as

`UnifiedTheory.traceMap_evenSign_cycle.`

Consequently only T^3 descends through the even-sign quotient by commuting with each flip, checked for flipAB as

`UnifiedTheory.traceMap_iterate_three_comm_flipAB`

and by the two cyclic variants. This is a finite algebraic invariant of the trace dynamics, not an analytic positivity or RH statement and not a consequence of periodicity alone.

The topology component checks hidden-direction rigidity. The theorem

`UnifiedTheory.Dynamics.hidden_rigidity`

states that a continuous map from a preconnected space into a totally disconnected hidden fiber is constant on that connected input. Its profinite specialization

`UnifiedTheory.Dynamics.profinite_fiber_rigid`

is theorem 20.12: profinite hidden fibers admit no nonconstant continuous path from a preconnected source. This formalizes the rigidity part of the solenoid intuition, not the full solenoid transport or pullback-line program.

2.11 Critical-Line Ledger and Completed Zeta

The phase reading is

`UnifiedTheory.phase,`

with scaling ledger

`UnifiedTheory.Lambda`

and normalized modulus

`UnifiedTheory.normModulus.`

For a nontrivial prime-axis state, the theorem

`UnifiedTheory.unitarity_line_iff`

states that

$$\operatorname{Re}(s) = 1/2 \iff \operatorname{normModulus}(s, a) = 1.$$

Equivalently, the normalized phase is unitary exactly on the critical line. The reflection

`UnifiedTheory.J`

is defined by $J(s) = 1 - \bar{s}$, and the checked theorem

`UnifiedTheory.J_fixed_iff`

states that its fixed line is exactly $\operatorname{Re}(s) = 1/2$. The scaling ledger changes sign under reflection:

`UnifiedTheory.Lambda_mirror.`

The sign criteria are also checked:

`UnifiedTheory.Lambda_pos_iff, UnifiedTheory.Lambda_neg_iff.`

For a nontrivial state, the local scaling ledger is positive exactly on the side $\operatorname{Re}(s) > 1/2$ and negative exactly on the side $\operatorname{Re}(s) < 1/2$. Finally,

`UnifiedTheory.ontic_balance_on_critical_line`

proves that a zero local scaling balance on a nontrivial state forces $\operatorname{Re}(s) = 1/2$. This is theorem 24.8 in its phase-ledger form: the implication from ontic balance to the critical line is unconditional.

The zeta bridge introduces

`UnifiedTheory.ProjectedZero`

for nontrivial zeros of `riemannZeta` and

`UnifiedTheory.OnticZero`

for projected zeros whose local scaling ledger is balanced on every nontrivial prime-axis state. The theorem

`UnifiedTheory.onticZero_re`

proves that every ontic zero lies on the critical line. The bridge theorem

`UnifiedTheory.rh_iff_all_projected_ontic`

states

`RiemannHypothesis` $\iff$ every projected zero is an ontic zero.

This is theorem 24.10 as a ledger restatement of `mathlib`'s `RiemannHypothesis`. It is not a proof of RH, because neither side of the equivalence is discharged unconditionally.

The ordinary zeta side contains the Euler product

`UnifiedTheory.zeta_eulerProduct,`

which states that for $\text{Re}(s) > 1$,

$$\prod_p' (1 - p^{-s})^{-1} = \zeta(s).$$

This is theorem 22.3: the analytic avatar of the free prime-axis monoid from theorem 4.4.

The no-zero-line facts are

`UnifiedTheory.zeta_ne_zero_of_one_le_re,` `UnifiedTheory.projectedZero_re_lt_one.`

They say that $\zeta(s) \neq 0$ when $\text{Re}(s) \geq 1$, and that every projected nontrivial zeta zero has $\text{Re}(s) < 1$. This tightens the zero ledger to the critical strip boundary; it does not place those zeros on the critical line.

The completed-zeta interface is also checked. The theorem

`UnifiedTheory.completed_functional_equation`

states

$$\Lambda(1 - s) = \Lambda(s)$$

for `mathlib`'s `completedRiemannZeta`. The explicit `riemannZeta` functional equation with Gamma and cosine factors is

`UnifiedTheory.riemannZeta_explicit_functional_equation.`

The zero-reflection theorem

`UnifiedTheory.completed_zero_reflect`

states that a zero of `completedRiemannZeta` at s gives a zero at $1 - s$. This is one half of the reflected-zero pairing. It gives symmetry about the critical line, not location on the critical line.

2.12 Weil Explicit-Formula Small-Support Reduction

We formalize the support-vanishing mechanism of the Weil explicit formula abstractly, not the full explicit formula. An explicit-formula functional is modeled by `UnifiedTheory.WeilFunctional`: an archimedean carrier together with the explicit prime-power side

$$\text{primeSide}(g) = \sum_p \sum_{k \geq 0} (\log p) (g((k+1) \log p) + g(-(k+1) \log p)),$$

which is `UnifiedTheory.primeSide`. When the test function g is supported in the open window $(-\log 2, \log 2)$, every prime-power frequency satisfies $(k+1)\log p \geq \log 2$ and so falls outside the window; hence the prime side vanishes term by term, `UnifiedTheory.primeSide_eq_zero_of_smallSupport` (unconditional). On such test functions the functional therefore reduces to its archimedean part, `UnifiedTheory.WeilFunctional.eval_smallSupport`, and is nonnegative whenever the archimedean part is, `UnifiedTheory.WeilFunctional.eval_nonneg_of_smallSupport`. The archimedean positivity is supplied as an explicit hypothesis: it is the Connes–Consani archimedean positivity, an open leaf that this development does not prove.

The reduction is non-vacuous: the single-point witness $g = \mathbf{1}_{\{\log 2\}}$ gives $\text{primeSide}(g) = \log 2$, checked as `UnifiedTheory.primeSide_witness` and `UnifiedTheory.primeSide_not_identically_zero`, so the prime side genuinely detects prime-power logarithms and $\log 2$ is the exact lower boundary of detectable frequencies.

Widening the window by one arithmetic threshold makes the finite-packet structure explicit. Since 2 is the only prime power below 3, a test function supported in $(-\log 3, \log 3)$ leaves exactly one surviving prime term:

$$\text{primeSide}(g) = \log 2 \left(g(\log 2) + g(-\log 2) \right),$$

checked as `UnifiedTheory.primeSide_eq_single_of_support_log3`. On this chamber the functional equals its archimedean part plus a single explicit prime shift, `UnifiedTheory.WeilFunctional.eval_log3_c`, the first nonzero finite prime side (the window $(-\log 2, \log 2)$ had none), and the semi-local finite-packet shape. Positivity is still not asserted: the archimedean term stays an open leaf.

The chamber structure extends to every window width. For a test function supported in the open window $(-\log N, \log N)$ with $N \geq 2$, every prime $p \geq N$ has smallest frequency $\log p \geq \log N$ outside the window, so it contributes nothing to the prime side, `UnifiedTheory.primeSide_inner_vanishes`. Consequently the outer sum collapses to a finite sum over the primes below N :

$$\text{primeSide}(g) = \sum_{\substack{p < N \\ p \text{ prime}}} \sum_{k \geq 0} (\log p) \left(g((k+1)\log p) + g(-(k+1)\log p) \right),$$

checked as `UnifiedTheory.primeSide_eq_sum_of_support`. On bounded support the Weil prime side is therefore a genuine finite object—the finite-packet cutoff—with the arithmetic thresholds $\log 2, \log 3, \log 4, \dots$ marking where each new prime power enters. This is the geometric prime side only; positivity is not asserted and the archimedean term remains an open leaf.

The prime-side frequencies connect the two layers. Under the canonical indexing $(p, k) \mapsto p^{k+1}$ of prime powers, the Weil frequency $\log(p^{k+1}) = (k+1)\log p$ carries the golden label $\Omega_\varphi(p^{k+1}) = S(k+1)$, checked as `UnifiedTheory.primePower_frequency_goldLabel`, so the explicit-formula and golden-weight layers meet on prime powers. Moreover the frequency map $(p, k) \mapsto (k+1)\log p$ is injective, checked as `UnifiedTheory.weilFreq_injective`: the prime side is indexed by distinct prime powers with no frequency collisions and no duplicate (p, k) data.

The count of surviving frequencies is a step function of the window. Writing $\text{ppChamber}(N) = \{q < N : q \text{ a prime power}\}$ for the prime powers below N , the set is monotone in N , `UnifiedTheory.ppChamber_monotone`, and its cardinality increases by exactly one precisely at prime-power boundaries, `UnifiedTheory.ppChamber_cardinality_gains`, and `UnifiedTheory.ppChamber_gains_iff`: the finite prime side acquires a new frequency exactly at each prime power, never between them and never two at once.

This formalizes the support-vanishing reduction only; unconditional Weil positivity, the explicit formula’s zero side, and RH are not claimed.

2.13 Concrete Archimedean Weil Functional

The archimedean term of the zeta explicit formula, in source ledger form 26.3/26.4, is made concrete in the Connes–Consani positivity sign. With $C_\infty = \gamma + \log(4\pi)$, `UnifiedTheory.archConst`,

and the desingularized kernel

$$K_g(t) = \frac{e^{t/2}(g(t) + g(-t)) - 2g(0)}{e^t - e^{-t}} \quad (t \neq 0), \quad K_g(0) := g(0)/2,$$

`UnifiedTheory.archKernel`, the compact-cutoff functional on support in $(-\log 2, \log 2)$ is

$$\text{archSmall}(g) = -C_\infty g(0) - \int_0^{\log 2} K_g(t) dt - g(0) \log \frac{e^{\log 2} - 1}{e^{\log 2} + 1},$$

`UnifiedTheory.archSmall_def`. The kernel’s singularity at the origin is removable: for g differentiable at 0, $K_g(t) \rightarrow g(0)/2$ as $t \rightarrow 0^+$, `UnifiedTheory.tendsto_archKernel_zero`. Pairing this with the standard weighted prime side $\sum_p \sum_{m \geq 1} \log p \cdot p^{-m/2} (g(m \log p) + g(-m \log p))$, `UnifiedTheory.primeSideZeta`, which also vanishes on small support, `UnifiedTheory.primeSideZeta_eq_zero_of`, realizes the kernel-internal form of the explicit formula: on small support the functional reduces to its archimedean part, `UnifiedTheory.concreteW_eq_arch_of_smallSupport`.

The well-definedness side is now complete: `UnifiedTheory.archKernel_intervalIntegrable` gives integrability on $[0, \log 2]$, `UnifiedTheory.archTail_log2` evaluates the cutoff tail as $-\log 3$, and `UnifiedTheory.convSquare_even` checks that convolution squares are even.

Positivity itself remains open. The target is registered precisely: the archimedean positivity on convolution squares, `UnifiedTheory.ArchimedeanPositivity`, is the Connes–Consani small-support archimedean positivity, source 26.9(iii), a month-scale analytic core that this development does not prove and that does *not* imply RH. The prime-side positivity limit is the RH wall. The equivalence of the spectral digamma form and the real-integral form, `UnifiedTheory.ArchSpectralMatchesReal`, is likewise registered as an open leaf, blocked on the digamma Gauss-integral representation. These are precise open propositions, not theorems; RH remains a ledger restatement.

2.14 Heart-Electrode Residual-Identity Foundation

The elementary closed foundation of the heart-chapter residual spectral identity is now separated from the analytic frontier. The audible-window Mellin transform is `UnifiedTheory.gTilde`; in source H.2 it is

$$\tilde{g}(s) = \frac{4 \sinh^2(s/2)}{s^2},$$

with removable center `UnifiedTheory.gTilde_zero`, $\tilde{g}(0) = 1$, and limit theorem `UnifiedTheory.tendsto_gTilde`. The window is moreover entire: `UnifiedTheory.gTilde_entire` proves $\tilde{g}$ analytic on all of $\mathbb{C}$, the removable singularity at $s = 0$ discharged by Riemann’s theorem because the punctured-neighbourhood limit equals the defined value 1, so the closed Mellin window carries no pole to obstruct the spectral sum. This is the closed Mellin-window part of the slogan “station two empty, station three fully audible.”

The audible window is also nonvanishing on the open right half-plane: `UnifiedTheory.gTilde_ne_zero_of_re` proves $\tilde{g}(s) \neq 0$ whenever $\text{Re } s > 0$, since the only zeros of $4 \sinh^2(s/2)/s^2$ are those of $\sinh(s/2)$ at $s \in 2\pi i\mathbb{Z}$ (the imaginary axis). This is the analytic core of the heart-native RH criterion of the source heart-electrode volume (H.7): because $\tilde{g}(\rho^*) \neq 0$ for any off-line zero $\rho^* = \beta + i\gamma^*$ with $\beta > 0$, some annular electrode in the family $\{g(t/\tau) R(e^{-t})\}$ hears it, and Landau’s oscillation theorem then forces the windowed sum to grow like $\delta^{-\beta}$. On the critical line the window decays: `UnifiedTheory.gTilde_critical_line_decay` proves $\|\tilde{g}(\frac{1}{2} + i\gamma)\| \leq 4 \cosh^2(\frac{1}{4})/\gamma^2$, an $O(\gamma^{-2})$ bound that makes the spectral (station-three) sum termwise controlled, via $\|\sinh(s/2)\| \leq \cosh(\text{Re}(s/2)) = \cosh(\frac{1}{4})$ (helper `UnifiedTheory.norm_sinh_le_cosh_re`) and $\|s\|^2 \geq (\text{Im } s)^2$. A subsequent literature-review round of the source volume corrects the lineage of the two headline statements, and we record that attribution honestly. The full criterion H.7 (RH $\iff$ every electrode sum is $O_\varepsilon(\delta^{-1/2-\varepsilon})$) is the spectral-window-family form of the Gerhold–Littlewood

criterion: Gerhold (2009, arXiv:0905.2070, Prop. 2) already gives $\text{RH} \iff \sum_n \mu(n)z^n = O((1-z)^{-1/2-\varepsilon})$, which is the Abel-window special case; and the unconditional heartbeat lower bound H.8 ($S_G(\delta) = \Omega(\delta^{-1/2})$) is a quantified window reproof of Delange (2000, Acta Arith. 92, 59–70), who proved the unconditional $\Omega_{\pm}((1-z)^{-1/2})$. What is genuinely new to the volume, and not a re-labelling of these classics, is the window-*family* refinement with per-window computable constants and, checked here, the “no-miss” mechanism: the right-half-plane non-vanishing `UnifiedTheory.gTilde_ne_zero_of_re_pos` is exactly what guarantees no off-line zero can hide from the electrode family, a step with no counterpart in the cited literature. The full criterion and lower bound themselves require Landau’s oscillation (Ω) theorem, which is not in `mathlib` (a candidate contribution in its own right), so they remain open leaves; the numerical certificate behind H.8, the double-contour value $\hat{R}(\rho_1) \approx 3.35 \times 10^{-8} \neq 0$, is empirical, and the “tax-reduction” claim H.9 (that the window family needs only a polynomial shadow of the Gonek–Hejhal bound $\sum_{\gamma \leq T} |\zeta'(\rho)|^{-2} \ll T$ for absolute convergence) depends on that zero-moment input and stays open. These sharpen the “modulo critical oscillation” qualifier on $D(1^-) = 2$ from a footnote toward a load-bearing statement: the constant 2 is the bone, the oscillation is the signal.

The concrete arithmetic left side of source H.1 is the windowed Möbius sum

$$S_g(N, \delta) = \sum_{k < N} \mu(k+1) g((k+1)\delta),$$

encoded as `UnifiedTheory.heartMobiusSum`. Its linearity in the window is checked by `UnifiedTheory.heartMobius` and `UnifiedTheory.heartMobiusSum_smul`. The analytic station-two input is `UnifiedTheory.one_div_riemannZet`, as $s \rightarrow 1$, $1/\zeta(s) \rightarrow 0$, so the zeta pole gives the reciprocal zeta factor a zero at the pole.

The source-6.147 one-page algebra of the $\bar{r}$ -chain is now closed in Lean. The heart δ -table term is $\delta_k = \text{fib}(k+1) \varphi^{-(k+2)}$, `UnifiedTheory.heartDelta`, with $\delta_0 = 1/\varphi^2$, `UnifiedTheory.heartDelta_zero`. Its geometric ψ -weighted sum has the closed value

$$\sum_k \psi^k \delta_k = \frac{1}{2\varphi}, \quad \bar{r} = -\frac{1}{2\varphi},$$

proven as `UnifiedTheory.rbar_series_eq` (with `UnifiedTheory.rbar_eq_neg`) by splitting the series through Binet into two geometric series, both of norm below one, the closing step being $\sqrt{5} = 2\varphi - 1$. The dual-seal constant is $C_0 = \sqrt{5}/2 - 1/(2\varphi) = \varphi/2$, `UnifiedTheory.heartC0_eq`, i.e. $C_0 = c^* \kappa + \bar{r}$ with $c^* = \sqrt{5}\varphi$ and $\kappa = 1/(2\varphi)$. These are the elementary constants of the heart chapter; they do not close $D(1^-) = 2$, whose full statement stays open below.

What is proven here is only this elementary closed foundation: the Mellin window, the windowed Möbius sum and its linearity, the zero of $1/\zeta$ at the pole, the δ -table constants $\bar{r}$ and C_0 , and the golden/fiber/Binet identities recorded above. The deep heart pieces are not closed algebra and are not formalized: the full residual spectral identity

$$S_g(\delta) = \text{station } 2 + \sum_{\rho} \frac{\tilde{g}(\rho) \delta^{-\rho}}{\zeta'(\rho)} + O(\delta^2),$$

the oscillation-growth statement H.4, the wall-coefficient identity H.6, and the shell positivity $F_d > 0$ on $(0, 1)$ all depend on the empirical Sturmian budget shell. That shell consists of frontier numerical facts: F_d is Sturmian-type, the deepest zero is $x_0 \approx 0.592093$, the oscillation is $B_n \sim 3.14(-1)^n x_0^{-n}$ with $R^2 \approx 1.0000$, the identity $F_d \equiv P$ holds to 80 digits, and the Lambert net form is

$$H = \sum_n \frac{B_n x^n}{1 - x^n}.$$

These are registered as precise open propositions, not theorems.

The heart main result $D(1^-) = 2$ holds modulo critical oscillation. This qualifier is a discovery about the structure of the problem: the residual sits in the same room as RH-type oscillation. It is not a dent in the problem itself, and it must not be dropped. RH remains a ledger restatement; the archimedean positivity leaf O-6 remains a month-scale open leaf.

2.15 Laplace-Transform Abscissa and Half-Plane Holomorphy

The analytic base of the Landau-oscillation engine of source F-7 is checked as reusable infrastructure, mirroring mathlib's Dirichlet-series `LSeries.abscissaOfAbsConv` with an integral in place of the sum. A `UnifiedTheory.LaplaceData` packages a real weight f on $[u_0, \infty)$ with its Laplace transform `UnifiedTheory.LaplaceData.F`, $F(s) = \int_{u \geq u_0} f(u) e^{-su}$, and the absolute-convergence predicate `UnifiedTheory.LaplaceData.Converges`. Convergence is monotone in the real part: `UnifiedTheory.LaplaceData.Converges.mono` shows that if the integral converges at σ it converges at every $\sigma' \geq \sigma$, because on $u \geq u_0$ one has $|f(u)e^{-\sigma'u}| \leq e^{-(\sigma'-\sigma)u_0} |f(u)e^{-\sigma u}|$. This yields a well-defined convergence abscissa `UnifiedTheory.LaplaceData.abscissa` $\sigma_c \in \mathbb{R}$ and half-plane convergence `UnifiedTheory.LaplaceData.converges_of_abscissa_lt`: $\sigma_c < \operatorname{Re} s$ forces the integrand to be integrable, so $F(s)$ is genuinely defined on $\{\operatorname{Re} s > \sigma_c\}$ (`UnifiedTheory.LaplaceData.integrableOnHalfPlane`).

On that half-plane F is holomorphic, with an explicit derivative: `UnifiedTheory.LaplaceData.hasDerivAt_F` proves, for $0 \leq u_0$ and $\sigma_c < \operatorname{Re} s_0$, that F is complex-differentiable at s_0 with $F'(s_0) = \int_{u \geq u_0} f(u)(-u)e^{-s_0 u}$, by differentiation under the integral sign (mathlib's `hasDerivAt_integral_of_dominated_1`). Holomorphy on the whole half-plane, `UnifiedTheory.LaplaceData.differentiableOn_F`, is its immediate corollary. The ∂_s -integrand $f(u)(-u)e^{-su}$ is itself integrable on the half-plane (`UnifiedTheory.LaplaceData.integrableOn_derivIntegrand`); the domination that validates the parametric derivative absorbs the polynomial factor $|u|$ through $u \leq \varepsilon^{-1}e^{\varepsilon u}$, uniformly on a neighbourhood $\{\sigma_1 < \operatorname{Re} s\}$ with $\sigma_c < \sigma_2 < \sigma_1 < \operatorname{Re} s_0$ and $\varepsilon = \sigma_1 - \sigma_2$ (extracted by `UnifiedTheory.LaplaceData.exists_real_between`). This is exactly the analytic prerequisite the oscillation argument of the heartbeat criterion would consume. The explicit form $F'(s) = \int (-u)f(u)e^{-su}$ is the $n = 1$ base of the iterated-derivative tower $F^{(n)}(s) = \int (-u)^n f(u)e^{-su}$ on which the Landau sign argument (non-negativity of f forcing $(-1)^n F^{(n)} \geq 0$) would run. The step from $F^{(n)}$ to $F^{(n+1)}$ needs no n -th-order parametric machinery: since $F^{(n)}$ is itself the transform of the polynomial-shifted weight $(-u)^n f(u)$, one derivative multiplies the weight by $-u$, and the shift `UnifiedTheory.LaplaceData.mulNegU` keeps the convergence abscissa inside the same half-plane (`UnifiedTheory.LaplaceData.abscissa_mulNegU_lt`: a polynomial factor does not move it, the $|u|$ being absorbed by the same $u \leq \varepsilon^{-1}e^{\varepsilon u}$ bound), a step made precise by `UnifiedTheory.LaplaceData.hasDerivAt_F'`: the first derivative of F is *literally* the transform of the shifted weight, $F' = (\text{mulNegU}).F$. Reapplying it at every order is now machine-checked: `UnifiedTheory.LaplaceData.abscissa_iterate_lt` keeps every shifted weight in the same half-plane, and the closed form $F^{(n)} = (\text{mulNegU}^{[n]}).F$ is proved as `UnifiedTheory.LaplaceData.iteratedDeriv_F`. The iterated derivative tower is complete at this analytic level; the remaining open leaves are the sign bookkeeping and the Landau Ω -theorem these transforms are meant to feed. The sign bookkeeping is now machine-checked: when $f \geq 0$, $(-1)^n F^{(n)}(\sigma) \geq 0$ on the real half-plane (`UnifiedTheory.LaplaceData.sign_control`); in complex-transform form this is `UnifiedTheory.LaplaceData.sign_control_iteratedDeriv_re`, using the real-axis identity `UnifiedTheory.LaplaceData.F_ofReal`. The reverse abscissa direction needed for the final contradiction is also checked: `UnifiedTheory.LaplaceData.abscissa_le_of_converges` proves that any convergent real exponent lies to the right of σ_c , and `UnifiedTheory.LaplaceData.not_converges` gives the strict-left nonconvergence corollary. The Landau contradiction itself remains open and is not registered as a Lean theorem; it still requires the analytic-continuation Taylor-radius argument and the Tonelli exchange turning the nonnegative Taylor series back into absolute convergence of the Laplace integral.

2.16 The φ -Coordinate System (Golden Coordinate Ring, L0)

The kernel's golden coordinate ring is given an axiomatic front as the base layer L0 of the φ -coordinate system (GCS). The carrier is $\mathcal{O} = \mathbb{Z}[\varphi] = \mathbb{Z} + \mathbb{Z}\varphi$ with $\varphi^2 = \varphi + 1$, encoded as `UnifiedTheory.PhiInt` on integer pairs (a, b) (this sidesteps the half-integer of $\mathbb{Z}[\sqrt{5}]$: $\mathcal{O}$ is the full ring of integers of $\mathbb{Q}(\sqrt{5})$, not $\mathbb{Z}[\sqrt{5}]$). L0 checks the four defining properties as theorems over that structure, not as Lean axioms. Ring axiom C1 is a `CommRing` instance on `UnifiedTheory.PhiInt`. The Galois conjugation $\sigma: a + b\varphi \mapsto (a + b) - b\varphi$ (axiom C2) is an involutive ring automorphism: `UnifiedTheory.PhiInt.conj_involutive`, packaged as `UnifiedTheory.PhiInt.conjHom` and `UnifiedTheory.PhiInt.conjAut`. The norm $N(a+b\varphi) = a^2 + ab - b^2$ (axiom C3), `UnifiedTheory.PhiInt.norm`, equals $x \sigma(x)$ (`UnifiedTheory.PhiInt.norm_eq_mul_conj`) and is multiplicative: `UnifiedTheory.PhiInt.norm_mul` proves $N(xy) = N(x)N(y)$. For the units (axiom C4), $N(\varphi) = -1$ (`UnifiedTheory.PhiInt.norm_phi`) and φ is a unit (`UnifiedTheory.PhiInt.isUnit_phi`, with inverse $\varphi - 1 = \langle -1, 1 \rangle$), and the powers carry $N(\varphi^n) = (-1)^n$ (`UnifiedTheory.PhiInt.norm_phiPow`). This is the constitutional first page of the coordinate ring. Two of the three coordinate axes of layer L1 are also checked: the multiplicative-degree axis `UnifiedTheory.PhiInt.degMul` lands the golden powers on the integer grid (`UnifiedTheory.PhiInt.degMul_phiPow` proves $\deg_{\times}(\varphi^n) = n$), and the Zeckendorf digit axis `UnifiedTheory.PhiInt.zeckDigits` round-trips through decoding (`UnifiedTheory.PhiInt.decode_zeckDigits`). The geometric axis is defined as the circle coordinate $\{n\varphi\}$ (`UnifiedTheory.PhiInt.geomCoord`), but its three-distance (three-gap) theorem is absent from mathlib and left an open leaf. The GCS prime-family arithmetic core is machine checked: for every prime $p \neq 2, 5$, 5 is a square modulo p exactly when $p \equiv \pm 1 \pmod{5}$, `UnifiedTheory.five_isSquare_iff`; this is the discriminant-five reason for mod-5 classes rather than Fermat mod-4 classes. The full prime-ideal splitting in $\mathbb{Z}[\varphi]$, class number $h = 1$, the discrete-depth metric, and the representation theorem for self-similar structures remain open leaves pending the ring-of-integers layer.

2.17 Kernel Components and Ledger States

The kernel aggregate is intentionally thin. The status enum

`UnifiedTheory.Kernel.LedgerStatus`

has four constructors:

`open, closed, tail, semantic.`

The component module supplies aliases such as

`UnifiedTheory.Kernel.History, UnifiedTheory.Kernel.ExponentState,`
`UnifiedTheory.Kernel.CodeRegister, UnifiedTheory.Kernel.decode, UnifiedTheory.Kernel.normalize.`

These names stabilize the formalized slots of $\mathcal{K}_{\text{PZG}}$. They do not constitute a theorem that every component of the full source-theory kernel has been built as one closed record.

Two ledger utilities are checked under `UnifiedTheory.Kernel`. The falsifiability theorem

`UnifiedTheory.Kernel.least_counterexample`

states that a false decidable universal claim over $\mathbb{N}$ has a least counterexample, with all smaller values satisfying the predicate. This is theorem 13.2 as a level-0 finite certificate. The tail-control predicate

`UnifiedTheory.Kernel.TailControlled`

is closed under addition by

`UnifiedTheory.Kernel.tailControlled_add,`

and a residual bounded by a tail budget tending to zero itself tends to zero by

`UnifiedTheory.Kernel.tendsto_zero_of_tailControlled.`

These are the checked tail-additivity and squeeze facts 12.4–12.5.

3 Generated Results (intuition→formalization loop)

The formalization cycle did more than ingest source-document theorems. It now has two idea sources: the local intuition→formalization loop, and a side-channel oracle loop in which ChatGPT Pro via omega-oracle proposes candidate statements. The oracle is only a conjecture source; a result enters this paper only after local Lean certification, so the workflow is multi-model assistance at the research-search layer rather than a replacement for kernel checking.

Own-intuition results

- (G1) **Sharp Beatty deficit range.** `UnifiedTheory.deficit_trichotomy_sharp` gives witnesses for all three values: $cDef(0,0) = 0$, $cDef(1,1) = 1$, and $cDef(2,2) = -1$. With `UnifiedTheory.cDef_mem`, the range is exactly $\{-1, 0, 1\}$.
- (G2) **Cassini–Fricke nondegeneracy.** `UnifiedTheory.cassini_fricke_ne_zero` proves that the invariant is nonzero exactly when $x \neq 0$ and $y \neq 0$, separating nontrivial Fibonacci-Hamiltonian parameters from the degenerate track $xy = 0$.
- (G3) **n -ary ledger composition.** `UnifiedTheory.Cn_append` states $C_n(as++bs) = C_n(as) + C_n(bs) + cDef(as.sum, bs.sum)$, the coboundary composition law for concatenated finite ledgers.
- (G4) **Gödel-code growth.** The theorem `UnifiedTheory.SelfCode.godelEncode_ge` proves $2^{\text{length}(l)} \leq \text{godelEncode}(l)$, an exponential lower bound for arithmeticized self-code registers.
- (G5) **Critical-side sign test.** The declarations

`UnifiedTheory.Lambda_pos_iff`
`UnifiedTheory.Lambda_neg_iff`

prove $0 < \Lambda_s(a) \iff \text{Re}(s) > 1/2$ and $\Lambda_s(a) < 0 \iff \text{Re}(s) < 1/2$ for nontrivial states, without asserting that zeta zeros lie on the line.

- (G6) **Critical reflection is involutive.** `UnifiedTheory.J_involutive` proves $J(J(s)) = s$ for the reflection $J(s) = 1 - \bar{s}$, giving the structural involution behind the zero-quartet symmetry ledger.
- (G7) **Normalized PZG addition laws.** The declarations

`UnifiedTheory.PZGTable.normAdd_comm`
`UnifiedTheory.PZGTable.normAdd_assoc`

certify commutativity and associativity through the decode equivalence, matching the free commutative prime-axis monoid behavior of theorem 4.4.

- (G8) **Phase readings never vanish.** `UnifiedTheory.phase_ne_zero` proves $\Phi_s(a) = \exp(-sL(a)) \neq 0$. In this formal layer, a zero is a projected cancellation phenomenon, not disappearance of the underlying phase vector.

- (G9) **Golden weight ranks the divisibility lattice.** Writing $\Omega_\phi(n) = \sum_{p|n} S(v_p(n))$ for the golden weight, the declarations

`UnifiedTheory.goldWeight_mono_of_dvd`
`UnifiedTheory.goldWeight_strictMono_of_dvd`

prove $m \mid n \Rightarrow \Omega_\phi(m) \leq \Omega_\phi(n)$ and, for a proper divisor $m \mid n$ with $m \neq n$ (and $n \neq 0$), the strict $\Omega_\phi(m) < \Omega_\phi(n)$. Since S is nonnegative and strictly monotone, golden weight is an order-embedding of the divisibility lattice $(\mathbb{N}^+, \mid)$ into $(\mathbb{Z}, \leq)$, strictly into $(\mathbb{Z}, <)$, matching the lattice reading of the fixed-length self-code below.

- (G10) **Golden weight is a quasimorphism.** With the cross-layer obstruction κ_ϕ , `UnifiedTheory.goldWeight_defect` proves $|\Omega_\phi(mn) - \Omega_\phi(m) - \Omega_\phi(n)| \leq |\text{supp}(m) \cap \text{supp}(n)|$: the additivity defect is bounded by the number of shared prime axes, via the deficit trichotomy $cDef \in \{-1, 0, 1\}$. Co-prime arguments give defect 0.

- (G11) **Golden-jump sequence is binary.** `UnifiedTheory.goldJump_mem` proves $\text{goldJump}(n) \in \{0, 1\}$: since $\tau = \varphi - 1 \in (0, 1)$, the floor difference $\lfloor (n+2)\tau \rfloor - \lfloor (n+1)\tau \rfloor$ is 0 or 1, so S grows by exactly 1 or 2 at each step—the golden Sturmian 0–1 word, companion to the window telescope (O2).

- (G12) **Shifted Zeckendorf–Beatty bridge is unconditional.** The declarations

`UnifiedTheory.minusWindow_holds`
`UnifiedTheory.shiftedZeckendorf_of_minusWindow`
`UnifiedTheory.S_eq_shifted_zeck`

discharge the former front-line bridge hypothesis: the asymmetric window is proved directly, then converted into the identity $S(v) = \sum_{i \in \text{encode } v} \text{fib}(i+1)$. This closes the roadmap P0 item without changing the analytic frontier.

- (G13) **Two-sided linear framing of the shift reading.** `UnifiedTheory.S_le_two_mul` proves $S(v) \leq 2v$ (from $\varphi < 2$) and `UnifiedTheory.succ_le_S` proves $v+1 \leq S(v)$ for $v \geq 1$ (from $\varphi > 3/2$), so $v+1 \leq S(v) \leq 2v$ is the arithmetic shadow of $1 < \varphi < 2$. Summing over prime axes, `UnifiedTheory.goldWeight_le_two_mul` and `UnifiedTheory.sum_succ_le_goldWeight` pin the golden weight between the total-with-multiplicity and total-plus-distinct prime counts: $\Omega + \omega \leq \Omega_\phi \leq 2\Omega$, where $\Omega = \sum_p v_p$ and ω is the number of distinct prime axes.

- (G14) **Algorithmic Zeckendorf normalization (theorem 5.7).** The multiset carry rewriting `UnifiedTheory.CarryFull` (rules $0 \mapsto \emptyset$, $1 \mapsto 2$, $\{i, i+1\} \mapsto \{i+2\}$, $\{i+2, i+2\} \mapsto \{i+3, i\}$) is value-preserving (`UnifiedTheory.CarryFull.fibVal_preserving`) and terminating (`UnifiedTheory.CarryFull.terminates`) via a weighted potential $\mu = \sum \text{ixWt}$ under which even the index-raising rule $1 \mapsto 2$ strictly decreases. It is confluent: `UnifiedTheory.exists_normal` reduces every multiset to a normal form and `UnifiedTheory.normal_form_unique` makes that form unique. Normal forms are exactly Zeckendorf sets (`UnifiedTheory.normal_isZeckMS`), so uniqueness follows from the uniqueness of Zeckendorf representations (`UnifiedTheory.normal_unique_of_same_value`) rather than critical-pair analysis: the rewriting strongly normalizes each multiset to the canonical Zeckendorf representation of its Fibonacci value.

- (G15) **Prime logarithms are $\mathbb{Q}$ -linearly independent.** The declaration

`UnifiedTheory.prime_log_linearIndependent`

proves that for a finite set of distinct primes and integer coefficients, $\sum_p k_p \log p = 0$ forces every $k_p = 0$: the relation becomes equality of two natural prime-power products,

and unique factorization reads off the exponents. This is the arithmetic root of the dense winding of the zeta phase line (the Brillouin picture of chapter 25)—the multiplicative independence of primes in logarithmic form.

Oracle-suggested results

(O1) **Beatty deficit as a τ -floor cocycle.** `UnifiedTheory.cDef_eq_floor_tau` proves $cDef(a, b) = \lfloor (a+1)\tau \rfloor + \lfloor (b+1)\tau \rfloor - \lfloor (a+b+1)\tau \rfloor$, with $\tau = \varphi - 1 = 1/\varphi$. Using $S(n) = n + \lfloor (n+1)\tau \rfloor$, the integer parts cancel and the deficit is controlled by the fractional rotation.

(O2) **Exact golden-jump window telescope.** `UnifiedTheory.goldJump_window_exact` proves that $goldJump(n) = S(n+1) - S(n) - 1 = \lfloor (n+2)\tau \rfloor - \lfloor (n+1)\tau \rfloor$ has length- N window sum $\lfloor (m+N+1)\tau \rfloor - \lfloor (m+1)\tau \rfloor$. This is the finite Sturmian balance-word telescope suggested by the oracle and then certified locally.

(O3) **Fixed-length Gödel coding is a lattice embedding.** For k distinct primes $p : \text{Fin } k \rightarrow \mathbb{N}$ and $Gvec(x) = \prod_i (p_i)^{x_i+1}$, the declarations

`UnifiedTheory.SelfCode.Gvec_dvd_iff`

`UnifiedTheory.SelfCode.Gvec_gcd` `UnifiedTheory.SelfCode.Gvec_lcm`

prove $Gvec(x) \mid Gvec(y) \iff \forall i, x_i \leq y_i$, and that gcd and lcm are computed coordinate-wise by min and max. The oracle proposed lifting injectivity of the fixed-length code to a full lattice embedding of $(\text{Fin } k \rightarrow \mathbb{N}, \leq)$ into $(\mathbb{N}^+, \mid)$; it was certified locally.

(O4) **Cross-layer golden-weight quasi-homomorphism.** With obstruction $\kappa_\phi(m, n) = \sum_p cDef(v_p(m), v_p(n))$, the declarations

`UnifiedTheory.goldWeight_mul`

`UnifiedTheory.goldWeight_mul_of_coprime`

prove $\Omega_\phi(mn) = \Omega_\phi(m) + \Omega_\phi(n) - \kappa_\phi(m, n)$, and that on coprime arguments the obstruction vanishes, so Ω_ϕ is strictly additive. This bridges the golden/deficit layer to the prime-factorization layer: the failure of Ω_ϕ to be a homomorphism is exactly the sum of pairwise Beatty deficits on shared prime axes. Proposed by the oracle and certified locally.

(O5) **Golden weight is a lattice valuation.** `UnifiedTheory.goldWeight_gcd_lcm_modular` proves $\Omega_\phi(\gcd(m, n)) + \Omega_\phi(\text{lcm}(m, n)) = \Omega_\phi(m) + \Omega_\phi(n)$, the exact modular law on the divisibility lattice. Where the product carries an obstruction κ_ϕ , meet and join carry none. Proposed by the oracle and certified locally.

(O6) **Golden weight of a Gödel code is coordinate-local.** `UnifiedTheory.SelfCode.goldWeight_Gvec` proves $\Omega_\phi(Gvec(P, x)) = \sum_i S(x_i + 1)$ for a fixed-length code over injective primes P : the golden rank depends only on the coordinate exponents, not on the concrete prime values. Proposed by the oracle and certified locally.

(O7) **Prime-axis edge jump is Sturmian.** `UnifiedTheory.goldWeight_mul_prime` proves $\Omega_\phi(pn) = \Omega_\phi(n) + 1 + goldJump(v_p(n))$, with `UnifiedTheory.goldWeight_mul_prime_pow` the p^N window form $\Omega_\phi(p^N n) - \Omega_\phi(n) = S(v_p(n) + N) - S(v_p(n))$. Multiplying by a prime raises golden weight by exactly 1 or 2, governed by the golden Sturmian word—sharpening strict divisor monotonicity. Proposed by the oracle and certified locally.

(O8) **Primitive two-square Zeckendorf atoms.** `UnifiedTheory.fib_odd_atom` proves that the odd Fibonacci number $\text{fib}(2k+3)$ is simultaneously a single-symbol Zeckendorf representation $[2k+3]$ and a primitive sum of two squares

$$\text{fib}(k+2)^2 + \text{fib}(k+1)^2$$

with consecutive Fibonacci numbers coprime. This gives an infinite certified family of atomic kernel objects with a quadratic witness. Proposed by the oracle and certified locally.

(O9) **Squarefree prime-axis Zeckendorf codes.**

For an injective prime stream p , `UnifiedTheory.AxisCode` sends A to

$$\prod_{i \in A} p_i.$$

Its prime factors are $A.\text{image } p$; it is injective and squarefree, with

`UnifiedTheory.axisCode_injective`
`UnifiedTheory.axisCode_squarefree`
`UnifiedTheory.axisCode_zeck_dvd_iff`

proving $p_i \mid \text{AxisCode } A \iff i \in A$. This is an arithmetic Gödel coding of Zeckendorf index sets as squarefree integers with no adjacent occupied prime axes. Proposed by the oracle and certified locally.

(O10) **The Fibonacci trace map is reversible and purely periodic.**

`UnifiedTheory.TraceStepEquiv` makes $(a, b, c) \mapsto (b, c, bc - a)$ a bijection with inverse $(a, b, c) \mapsto (ab - c, a, b)$. Hence

`UnifiedTheory.traceStep_zmod_pure_periodic`

shows that modulo every q , each state is purely periodic: there are no preperiodic tails and no state merging. The trace dynamics loses no information. Proposed by the oracle and certified locally.

(O11) **Fricke trace complement.** The theorem

`UnifiedTheory.trace_mul_add_trace_adjugate_mul`

proves

$$\text{tr}(AB) + \text{tr}(\text{adj}(A) B) = \text{tr } A \cdot \text{tr } B$$

for 2×2 matrices, and `UnifiedTheory.SL2_trace_inv_mul` gives

$$\text{tr}(A^{-1}B) = \text{tr } A \cdot \text{tr } B - \text{tr}(AB)$$

in SL_2 . Thus three trace coordinates determine the fourth. Proposed by the oracle and certified locally.

(O12) **Gödel-code divisor count.** The declarations

`UnifiedTheory.SelfCode.Gvec_divisors_card`
`UnifiedTheory.SelfCode.Gvec_primeFactors`

prove that the fixed-length Gödel code

$$\text{Gvec } p x = \prod_i p_i^{x_i+1}$$

has exactly $\prod_i (x_i + 2)$ divisors and pins its prime factors to $\text{image } p$. Divisors are exactly bounded coordinatewise exponent choices, with no off-axis divisors. Proposed by the oracle and certified locally.

- (O13) **Gödel codes have no infinite antichain (Dickson).** `UnifiedTheory.SelfCode.finiteAxis_dickson` shows $\text{Fin } k \rightarrow \mathbb{N}$ is well-quasi-ordered, and `UnifiedTheory.SelfCode.Gvec_fixedAxis_dividing_pair` transports this through the divisibility embedding: in any infinite sequence of fixed k -axis Gödel codes, some earlier code divides a later one. There is no infinite divisibility antichain inside a fixed prime-axis block.
- (O14) **Bounded Zeckendorf words have Fibonacci cardinality.** `UnifiedTheory.card_zeckPrefix` proves that the Zeckendorf words with all indices below $n+2$ number exactly $\text{fib}(n+2)$, via an explicit equivalence to $\text{Fin}(\text{fib}(n+2))$ (`UnifiedTheory.zeckPrefixEquivFin`); the key lemma `UnifiedTheory.zeck_sum_lt_fib_iff_bounded` identifies the initial interval $[0, \text{fib}(n+2))$ with the naturals whose Zeckendorf support is bounded by $n+1$: a finite counting theorem, not another uniqueness statement.

4 Scope and Honesty

The checked column has precise boundaries. The Zeckendorf shrinking-window and Beatty-floor routes are

`UnifiedTheory.deficitTrichotomy_holds`
`UnifiedTheory.cDef_mem.`

They are independent mechanisms; the bridge

`UnifiedTheory.S_eq_shifted_zeck`

is no longer a checked theorem under a separate bridge hypothesis. The assumption

`UnifiedTheory.ShiftedZeckendorfBeattyBridge.`

has been discharged by `UnifiedTheory.minusWindow_holds`, so the Zeckendorf shrinking-window and Beatty-floor routes now meet in one unconditional theorem. At the project-ledger level, the front-line conditional count is now zero; this does not promote any analytic outlook item into the checked column. The n -ary theorem `UnifiedTheory.Cn_abs_le` is the symmetric estimate $|C_n| \leq m - 1$; sharper asymmetric endpoint counts are not claimed. The normalization theorem

`UnifiedTheory.PZGTable.decode_normAdd`

is canonical normalization through the PZG bijection, not termination of a concrete carry-rewrite algorithm.

The zeta results include

`UnifiedTheory.zeta_eulerProduct`
`UnifiedTheory.zeta_ne_zero_of_one_le_re`
`UnifiedTheory.projectedZero_re_lt_one,`

along with completed functional equations and reflected completed-zeta zeros. These facts give prime-axis analyticity in $\text{Re}(s) > 1$, symmetry about the critical line, and exclusion of projected nontrivial zeros from $\text{Re}(s) \geq 1$. They do not prove RH.

The golden-weight layer contains the finite Euler factorization

`UnifiedTheory.goldGenAF_euler.`

This is an algebraic product over the finite set of prime axes dividing one natural number. It is not a formalization of the analytic Dirichlet series $\sum_n x^{\Omega_\varphi(n)} n^{-s}$, its convergence domain, continuation, or any ζ -relation.

The explicit-formula layer contains the small-support prime-side vanishing theorem

`UnifiedTheory.primeSide_eq_zero_of_smallSupport`

and the corresponding reduction and conditional nonnegativity declarations

`UnifiedTheory.WeilFunctional.eval_smallSupport`

`UnifiedTheory.WeilFunctional.eval_nonneg_of_smallSupport.`

The last statement assumes archimedean positivity; it does not prove the Connes–Consani archimedean positivity leaf, unconditional small-support positivity, the zero side of the explicit formula, or RH.

The declaration

`UnifiedTheory.rh_iff_all_projected_ontic`

is an equivalence with `RiemannHypothesis`. The theorem `UnifiedTheory.onticZero_re` proves that an already ontic zero lies on the critical line. The missing RH-strength step is still the universal upgrade from projected zeros to ontic zeros; there is no unconditional declaration of type `RiemannHypothesis`.

The checked self-code layer contains Cantor barriers, Kleene fixed points, and the Gödel sequence encoder. It still does not provide a verified instruction set, substitution evaluator, same-layer closure-decider contradiction, or full internalization tower. The checked dynamics layer contains cost growth and profinite hidden-fiber rigidity, not full solenoid transport.

5 Outlook

All items in this section are outside the present checked-result column and are **not yet formalized**. The analytic program still lacks a PZG heat-trace derivation, Mellin/theta bridge, the full Weil explicit formula with its zero side, tail-budget analysis, and a positivity or spectral argument upgrading projected zeros to ontic zeros. The checked Weil contribution is the support-vanishing reduction of the prime side under a small-support condition; Weil archimedean positivity, unconditional small-support positivity, and RH remain open leaves. The analytic Dirichlet series attached to the golden weight, including convergence, continuation, and any relation to ζ , also remains an open leaf. RH itself, O-5, and O-6 remain open leaves.

The source theory’s quasicrystal zeta object Z_{qc} , its Euler product or continuation theory, natural-boundary claims, diagonal feedback program, and pullback-line claim 6.32 remain outlook. The finite trace-map declarations

`UnifiedTheory.cassini_fricke`

`UnifiedTheory.cassini_fricke_ne_zero`

supply algebraic input for such a program, but they do not build Z_{qc} .

The intended research workflow is correspondence search followed by theorem extraction: prime-axis exponent vectors with Euler products, Beatty–Wythoff and Zeckendorf coding, $\mathbb{Q}(i)$ two-square norms with $\mathbb{Q}(\sqrt{5})$ deficit traces, Hecke–Mahler series with golden word factors, and hidden profinite fibers with rigidity. This manuscript reports only the correspondences backed by actual Lean declarations.

6 Reproducibility

The checked source lives under

`papers/unified_theory/lean4/UnifiedTheory.`

It is a mathlib project, separate from the mathlib-free BEDC core. The full verification command for the Lean project is

```
cd papers/unified_theory/lean4 && lake build.
```

The branch described here is intended to have that check green, with no `sorry` and no `axiom` in the `UnifiedTheory` sources; individual claims reduce to Lean type-checking of the declarations named above.